

When Sample Selection Bias Precipitates Model Collapse

Xinbao Qiao, Xianglong Du, Wei Liu, Jingqi Zhang, Peihua Mai, Meng Zhang, Yan Pang

Introduction

Web-scale corpora increasingly mix human data with model-generated content. Recursive training on synthetic data can narrow coverage and trigger model collapse (Shumailov et al., Nature 2024). Verification and sample selection are natural safeguards, but local validation may turn filtering into selection bias.

Theory

Proposition 1: Replace paradigm (Shumailov et al., Nature 2024)

Under the *Replace* paradigm, as $t \rightarrow \infty$,

$$\Sigma_t \xrightarrow{a.s.} 0, \quad \mathbb{E}[\mathbb{W}_2^2(\mathcal{N}_t, \mathcal{N}_0)] \rightarrow \infty.$$

Replacing data eventually removes variance and drifts away.

Proposition 2: Accumulate paradigm (Kazdan et al., ICML 2025)

Under unbiased accumulation,

$$\mathbb{E}[\bar{\Sigma}_t] \rightarrow \alpha_n \bar{\Sigma}_0.$$

$$\text{Thus } \mathbb{E}[\mathbb{W}_2^2] \rightarrow 2(1 - \sqrt{\alpha_n}) \text{Tr}(\bar{\Sigma}_0).$$

Unbiased accumulation keeps diversity from fully collapsing.

Theorem 1: local selection collapses diversity

With top- α selection toward a local ideal $\mathbf{u}^*$,

$$\|\bar{\mu}_t - \mathbf{u}^*\|^2 \rightarrow 0, \quad \bar{\Sigma}_t \rightarrow 0.$$

Local fidelity erases global coverage.

Local selection gains fit by erasing global coverage.

Theorem 2: diversity decays by power law

If the trajectory remains in the local basin,

$$\text{Tr}(\bar{\Sigma}_t) = \mathcal{O}_{a.s.}(t^{-\psi}).$$

Tail erosion has an explicit rate.

Once trapped locally, diversity shrinks at a predictable rate.

Theorem 3: drift controls true-manifold risk

For filtered distribution $\mathcal{D}_t$ and true manifold $\mathcal{D}^*$,

$$\mathcal{R}_{\mathcal{D}^*}(h_t; g^*) \leq 2\ell\epsilon \mathbb{W}_p(\mathcal{D}_t, \mathcal{D}^*) + \mathcal{R}_{\mathcal{D}_t}(h_t; g_t) + \mathcal{O}(\ell\delta).$$

Distribution drift upper-bounds risk on true data.

For more methodology and theoretical details, please refer to the full paper.



Synthetic data loops can fail when filters favor one region rare cases disappear.




Scan

OpenReview paper
and supplement.


Contact: xinbaoqiao@cuhk.edu.hk

Multivariate Gaussian Modeling


Gaussian




Anisotropic Gaussian




Gaussian mixture



Laplace



Collapse Signal



Airplane-only verification drives CIFAR-10 proportions toward Airplane, exposing systematic tail-mode pruning.

Experimental Synthesis

Across CIFAR-10, STL-10, and CelebA, skewed local references can make selection weaker than random sampling; broader distributional evidence improves FID, precision, and recall.

References

- [1] Shumailov et al. AI models collapse when trained on recursively generated data. Nature, 2024.
- [2] Kazdan et al. Collapse or Thrive: Perils and Promises of Synthetic Data in a Self-Generating World. ICML, 2025.

For additional experimental results, please refer to the full paper.

